

Dimension Reduction

Principal Components Analysis (PCA)

Dimension Reduction



Decrease the number of the variables (dimensions)

$$\mathbf{X} = \begin{pmatrix} x_1^{(1)} & x_1^{(2)} & \dots & x_1^{(p)} \\ x_2^{(1)} & x_2^{(2)} & \dots & x_2^{(p)} \\ \dots & \dots & \dots & \dots \\ x_n^{(1)} & x_n^{(2)} & \dots & x_n^{(p)} \end{pmatrix} \rightarrow \mathbf{Z} = \begin{pmatrix} z_1^{(1)} & z_1^{(2)} & \dots & z_1^{(p')} \\ z_2^{(1)} & z_2^{(2)} & \dots & z_2^{(p')} \\ \dots & \dots & \dots & \dots \\ z_n^{(1)} & z_n^{(2)} & \dots & z_n^{(p')} \end{pmatrix}$$

$$p' \ll p$$

Why?

- ▶ Find latent/hidden variables that are not/cannot be directly measured
- ▶ Reveal the hidden structure of the data
- ▶ Transform the feature space into variables that are not correlated, thus new variables can be used in different machine learning techniques

Principal Components Analysis (PCA)



Principal Components Analysis (PCA)

- ▶ PCA is a technique that can be used to simplify a dataset.
- ▶ It is a linear transformation that chooses a new coordinate system for the data set such that greatest variance by any projection of the data set comes to lie on the first axis (called the first principal component), the second greatest variance on the second axis, and so on.
- ▶ PCA can be used for reducing dimensionality by eliminating the later principal components.

Principal Components Analysis (PCA)

Each new Component/factor is a linear combination of all variables.

Imagine we have p variables ($X_j, j = 1, 2, \dots, p$)

$$Z_i = \sum_{j=1}^p w_{ij} X_j$$

$$Z_i = w_{i1}X_1 + w_{i2}X_2 + w_{i3}X_3 + \dots + w_{ip}X_p$$

where each pair of Z 's are orthogonal (correlation = 0)!

After estimating weights w , Z 's are ordered by their variance, with Z_1 having the largest variance and Z_p having the smallest variance.

Estimating principal components weights

$$Z_i = w_i' X$$


where w_k' is *eigenvector* of Σ (covariance matrix of initial variables $X_j, j = 1, 2, \dots, p$) corresponding to its k th largest *eigenvalue* λ_k .

Furthermore, if w_k is chosen to have unit length ($w_k' w_k = 1$), then $\text{var}(Z_k) = \lambda_k$.

A standard approach to estimate the unknown weights is the technique of Lagrange multipliers.

$$\lambda_1 \geq \lambda_2 \geq \lambda_3 \geq \dots \geq \lambda_p$$

Geometric picture of principal components



- ▶ The 1st PC Z_1 is a minimum distance fit to a line in X space
- ▶ The 2nd PC Z_2 is a minimum distance fit to a line in the plane perpendicular to the 1st PC

PCA in R: Example



Annual food consumption in Armenia

```
consumption<-read.csv("consumption.csv")  
str(consumption)
```

```
## 'data.frame':    4000 obs. of  6 variables:  
## $ bread_1_flour_white.kg: num  100 133.89 0.5 1.25 50 ...  
## $ bread_2_cereal.kg    : num   3.33 1.67 1.5 7.5 1.25 25 1.25 2.5 1.5 3 ...  
## $ bread_3_rice.kg      : num   1.67 1.33 1.5 3.75 2.5 10 1.25 1.67 1.5 3.5  
## $ bread_4_beans.kg     : num   1.33 1 0.5 1.25 2.5 5 0.75 1 2.5 2 ...  
## $ bread_5_macaroni.kg  : num    6 6.67 5 3.75 7.5 30 2.5 5 10 8.75 ...  
## $ bread_6_lavash.kg    : num  155 166.7 50 195 58.8 ...
```


Descriptive Analysis

```
round(data.frame(
  mean=sapply(consumption, mean),
  sd=sapply(consumption, sd),
  min=sapply(consumption, min),
  max=sapply(consumption, max),
  median=sapply(consumption, median)),3)
```

##		mean	sd	min	max	median
##	bread_1_flour_white.kg	38.569	70.198	0	550.00	5
##	bread_2_cereal.kg	3.016	5.182	0	187.50	2
##	bread_3_rice.kg	3.682	2.911	0	30.00	3
##	bread_4_beans.kg	2.772	2.602	0	30.00	2
##	bread_5_macaroni.kg	7.353	5.932	0	106.67	6
##	bread_6_lavash.kg	115.343	79.294	0	636.00	100

Let's take two products: rice and macaroni


```
ggplot(consumption, aes(x=bread_3_rice.kg, y=bread_5_macaroni.kg)) +  
  geom_point(color="blue", size=1.5)+xlim(0,30)+ylim(-5,100)
```



```
cor(consumption$bread_3_rice.kg, consumption$bread_5_macaroni.kg)
```

```
## [1] 0.603148
```

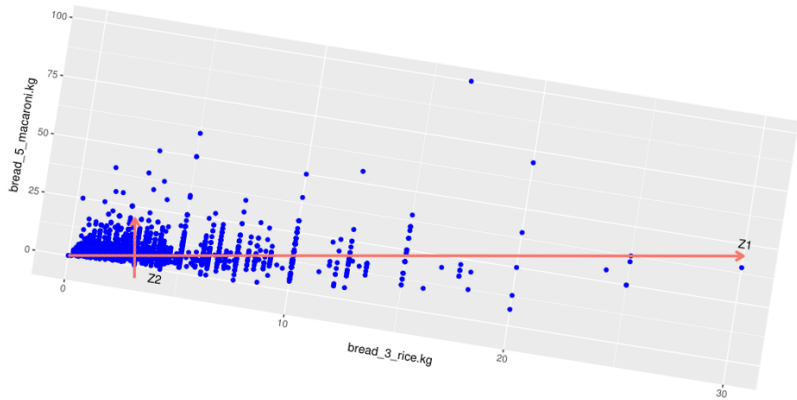
Let's take two products: rice and macaroni



- ▶ Z_1 accounts for the highest variance
- ▶ Z_1 and Z_2 are orthogonal, so the correlation is zero

This approach is called varimax orthogonal rotation, meaning that factors are orthogonal.

Rotated Space



Other considerations

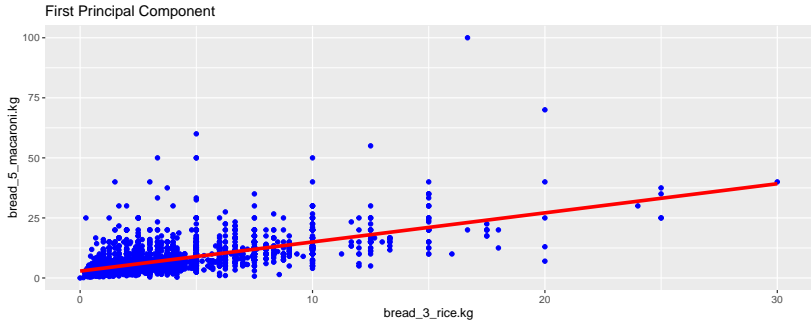
- ▶ Principal component analysis is otherwise called Factor Analysis
- ▶ Is done only with numeric variables
- ▶ PCA for categorical variables is called Correspondence analysis

Base R uses 2 methods of PCA,

- ▶ *princomp* uses covariance matrix
- ▶ *prcomp* uses Singular Value Decomposition
- ▶ The results are usually very similar

First Principal Component

```
ggplot(consumption, aes(x=bread_3_rice.kg, y=bread_5_macaroni.kg)) +  
  geom_point(color="blue", size=1.5)+xlim(0,30)+ylim(0,100)+  
  geom_smooth(method = "lm", se=F, size=1.5, col="red")+  
  ggtitle("First Principal Component")
```



Run PCA using base R **prcomp** function

```
p_comp <- prcomp(consumption[,c("bread_3_rice.kg", "bread_5_macaroni.kg")])
names(p_comp)
```

```
## [1] "sdev"      "rotation" "center"    "scale"     "x"
```

Principal components score for each observation

```
head(p_comp$x)
```

```
##           PC1      PC2
## [1,]  1.9338107  1.463289
## [2,]  1.4107266  2.002625
## [3,]  2.9347612  1.298978
## [4,]  3.3854453 -1.235166
## [5,]  0.2454652  1.165985
## [6,] -23.4705820  1.387580
```

Summary

```
summary(p_comp)

## Importance of components:
##               PC1      PC2
## Standard deviation    6.2264 2.2119
## Proportion of Variance 0.8879 0.1121
## Cumulative Proportion 0.8879 1.0000
```

Proportion of variance explained by each factor/component

- ▶ First component (PC1) explains 89% of the variance
- ▶ Second component explains 11% of the variance


How to compute proportions of explained variance with **standard deviations**

```
(6.2264)^2/((6.2264)^2+(2.2119)^2)
```

```
## [1] 0.8879423
```

Plot components

```
ggplot(as.data.frame(p_comp$x), aes(x = PC1, y = PC2)) +  
  geom_point(color = "blue", size = 1.5) + xlim(-5, 10) + ylim(-5, 5) +  
  theme(axis.title = element_text(size = 20))
```



Calculate PCA scores by hand

Lets Create a dataframe with original variables and the PCA

```
d1<-data.frame(consumption[,c("bread_3_rice.kg", "bread_5_macaroni.kg")], p_com)
summary(d1)
```

##	bread_3_rice.kg	bread_5_macaroni.kg	PC1	PC2
## Min.	: 0.000	Min. : 0.000	Min. : -95.976	Min. : -15.5461
## 1st Qu.:	1.670	1st Qu.: 3.750	1st Qu.: -2.281	1st Qu.: -0.8732
## Median :	3.000	Median : 6.000	Median : 1.501	Median : 0.2868
## Mean :	3.682	Mean : 7.353	Mean : 0.000	Mean : 0.0000
## 3rd Qu.:	5.000	3rd Qu.: 10.000	3rd Qu.: 3.882	3rd Qu.: 1.1382
## Max. :	30.000	Max. : 106.670	Max. : 8.151	Max. : 26.3113

You can see that the means for components is equal to zero (because of scalling)

Now lets claculate the PC scores manually

Mean for rice = 3.682

Mean for macaroni = 7.353

Rotation provides loadings

```
p_comp$rotation
```

```
##                PC1        PC2
## bread_3_rice.kg   -0.3250778 -0.9456873
## bread_5_macaroni.kg -0.9456873  0.3250778
```

```
head(d1, n=2)
```

```
##  bread_3_rice.kg bread_5_macaroni.kg      PC1      PC2
## 1              1.67                6.00 1.933811 1.463289
## 2              1.33                6.67 1.410727 2.002625
```

Lets calculate the score of the PC1 for the first case and PC2 for teh second case

```
(1.67-3.682)*(-0.3250778) + (6.00-7.353)*(-0.9456873)
```

```
## [1] 1.933571
```

```
(1.33-3.682)*(-0.9456873) + (6.67-7.353)*0.3250778
```

```
## [1] 2.002228
```

Correlations

```
round(cor(d1),4)
```

##	bread_3_rice.kg	bread_5_macaroni.kg	PC1	PC2
## bread_3_rice.kg	1.0000	0.6031	-0.6954	-0.7186
## bread_5_macaroni.kg	0.6031	1.0000	-0.9926	0.1212
## PC1	-0.6954	-0.9926	1.0000	0.0000
## PC2	-0.7186	0.1212	0.0000	1.0000

Lets go back to variance

Calculate total variance of the original data

```
cov(d1[,1:2])
```

##	bread_3_rice.kg	bread_5_macaroni.kg
## bread_3_rice.kg	8.47237	10.41425
## bread_5_macaroni.kg	10.41425	35.18871

Variance for rice = 8.47237; Variance for macaroni = 35.18871

Covariance = 10.41425

Total Variance = $8.47237 + 35.18871 = 43.66108$

% of variance explained by macaroni $35.18871 / 43.66108 = 80\%$

If we would like to do dimensionality reduction and keep only macaroni (highest variance) we would be able to keep only 80% of original information (lose 20%)

If we take PC1 we will be able to keep 89% of the original information

Run PCA for all bread products

Note: scale the initial variable to have a **unit variance**

```
colnames(consumption)
```

```
## [1] "bread_1_flour_white.kg" "bread_2_cereal.kg"
## [3] "bread_3_rice.kg"        "bread_4_beans.kg"
## [5] "bread_5_macaroni.kg"    "bread_6_lavash.kg"
```

```
p_comp <- prcomp(consumption, scale=T)
summary(p_comp)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6
## Standard deviation	1.638	1.0285	0.9132	0.7340	0.70605	0.62345
## Proportion of Variance	0.447	0.1763	0.1390	0.0898	0.08308	0.06478
## Cumulative Proportion	0.447	0.6233	0.7623	0.8521	0.93522	1.00000

By default the number of extracted components is equal to the number of variables in the model

- ▶ Now, how many components to take? Remember we are doing dimensionality reduction
- ▶ There are extracted as many components as many variables we have.
- ▶ When standardized, each variable has a variance/stdev of 1
- ▶ If the component has variance < 1 , then the component explains less than 1 variable from original dataset
- ▶ So take only components that have variance (also called eigenvalues) of greater than 1.
- ▶ Use Screeplot

```
screeplot(p_comp, type="l")
```




```
# Eigenvalues
p_comp$sdev^2
```

```
## [1] 2.6821930 1.0578477 0.8339716 0.5387955 0.4985022 0.3886898
```

First 3 components together explain the biggest part of the variance (67%), we can keep them, however eigenvalues suggest that 2 components are enough (eigenvalues > 1)

Percentage of explained variances

`fviz_eig(p_comp)`



Correlation matrix

```
df <- data.frame(consumption, p_comp$x[, 1:2])
cor_mat<-cor(df)
```

Lets get the correlation matrix in a way that by columns we will have PC components only and by the rows we will get the correlations

```
cor_mat <- cor_mat[! rownames(cor_mat) %in% c("PC1", "PC2"),
                  colnames(cor_mat) %in% c("PC1", "PC2")]
print(cor_mat, digits=3)
```

##	PC1	PC2
## bread_1_flour_white.kg	-0.385	0.8332
## bread_2_cereal.kg	-0.496	0.0870
## bread_3_rice.kg	-0.823	-0.0540
## bread_4_beans.kg	-0.763	0.1348
## bread_5_macaroni.kg	-0.800	-0.0825
## bread_6_lavash.kg	-0.623	-0.5728

```
print(cor_mat, digits=3)
```

##	PC1	PC2
## bread_1_flour_white.kg	-0.385	0.8332
## bread_2_cereal.kg	-0.496	0.0870
## bread_3_rice.kg	-0.823	-0.0540
## bread_4_beans.kg	-0.763	0.1348
## bread_5_macaroni.kg	-0.800	-0.0825
## bread_6_lavash.kg	-0.623	-0.5728

Lets see which variables are highly correlated with which factors, lets take a threshold of 0.6 for correlation coefficient


PC1 is described by rice, macaroni and beans

PC2 is described by flour_white

Visualization of the variables on the factor map


Correlation circle can help to visualize the most correlated variables

```
fviz_pca_var(p_comp, repel = TRUE) # Avoid text overlapping
```



Contributions of variables on PC1

```
fviz_contrib(p_comp, choice = "var", axes = 1)
```



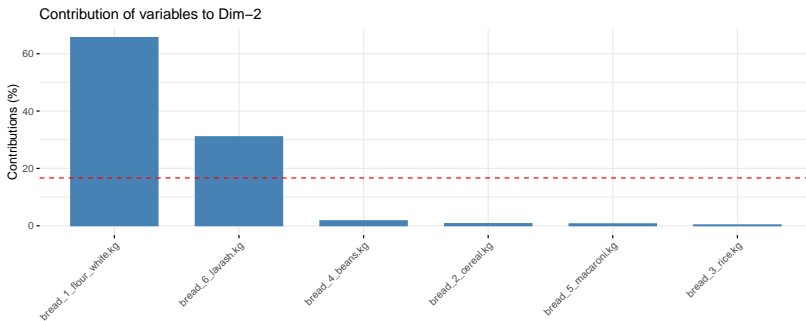
What means the red line on the graph?

If the contribution of the variables were uniform, the expected value would be $1/\text{length}(\text{variables}) = 1/6 = 16.67\%$.

The red dashed line on the graph above indicates the expected average contribution. For a given component, a variable with a contribution larger than this cutoff could be considered as important in contributing to the component.

Contributions of variables on PC2

```
fviz_contrib(p_comp, choice = "var", axes = 2)
```




Total contribution on PC1 and PC2

```
fviz_pca_contrib(p_comp, choice = "var", axes = 1:2)
```


```
## Warning in fviz_pca_contrib(p_comp, choice = "var", axes = 1:2): The
## function fviz_pca_contrib() is deprecated. Please use the function
## fviz_contrib() which can handle outputs of PCA, CA and MCA functions.
```

Contribution of variables to Dim-1-2



Control variable colors using their contributions

```
fviz_pca_var(p_comp, col.var="contrib")
```

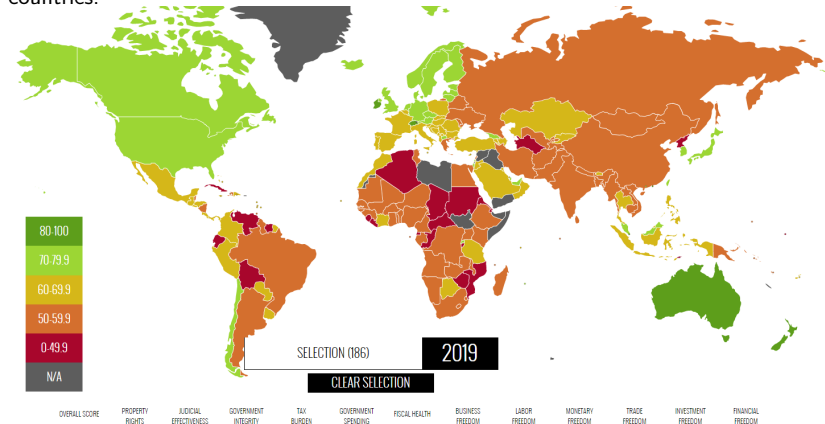


Construct Economic Freedom Index for countries with PCA



The Heritage Foundation measures an Economic Freedom Index of countries

The Index covers 12 freedoms – from property rights to financial freedom – in 186 countries.



Construct an Economic Freedom Index with PCA based on 6 freedom indices:

- ▶ Business.Freedom
- ▶ Labor.Freedom
- ▶ Monetary.Freedom
- ▶ Trade.Freedom
- ▶ Investment.Freedom
- ▶ Financial.Freedom

For such problems only first principal component could be used (ignoring other components)

```
df <- read.csv("Countries.csv")
df1 <- data.frame(Country = df[,2], df[,grepl(".Freedom", colnames(df))])
df1 <- df1[complete.cases(df1), ]
rownames(df1) <- df1$Country
head(df1)
```

##	Country	Business.Freedom	Labor.Freedom	Monetary.Freedom
## Afghanistan	Afghanistan	54.2	59.9	69.3
## Albania	Albania	79.3	50.7	81.4
## Algeria	Algeria	62.1	49.5	67.0
## Angola	Angola	58.5	40.4	70.6
## Argentina	Argentina	57.3	46.1	50.9
## Armenia	Armenia	78.5	72.4	72.8
##	Trade.Freedom	Investment.Freedom	Financial.Freedom	
## Afghanistan	66.0	0	0	
## Albania	87.7	70	70	
## Algeria	63.3	35	30	
## Angola	56.7	30	40	
## Argentina	66.7	50	50	
## Armenia	80.2	80	70	


```
p_comp <- prcomp(df1[, -1], scale. = T)
summary(p_comp)
```

```
## Importance of components:
```

##	PC1	PC2	PC3	PC4	PC5	PC6
## Standard deviation	1.8914	0.9440	0.71885	0.69545	0.58525	0.43396
## Proportion of Variance	0.5963	0.1485	0.08613	0.08061	0.05709	0.03139
## Cumulative Proportion	0.5963	0.7448	0.83092	0.91153	0.96861	1.00000


First principal component explain 60% of total variation!

```
screepplot(p_comp, type="l")
```



Visualization of the individuals using factoextra package

```
fviz_pca_ind(p_comp)
```



Run PCA one more time and extract only the first component

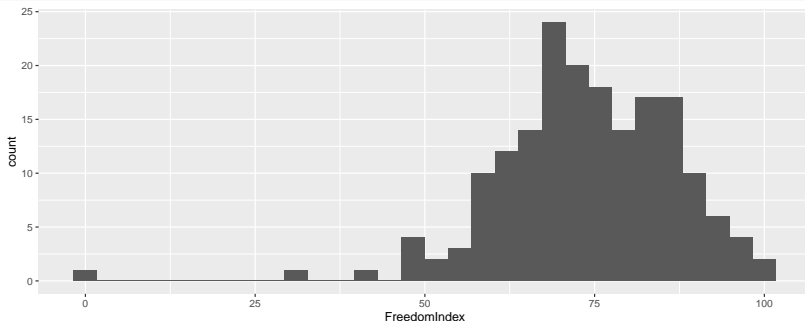
```
p_comp<-prcomp(df1[, -1], scale. = T, rank. = 1)
head(p_comp$x)
```

```
##                PC1
## Afghanistan -3.088722
## Albania      1.592204
## Algeria      -1.922102
## Angola       -2.169042
## Argentina    -1.872372
## Armenia      1.598681
```

Construct Economic Freedom Index with Min-Max normalization

```
df1$Score <- as.numeric(p_comp$x)
df1$FreedomIndex <- (df1$Score - min(df1$Score))/(max(df1$Score)-min(df1$Score))

ggplot(df1, aes(x = FreedomIndex)) + geom_histogram()
```



Correlation of Index and individual components

```
cor <- cor(df1[,c(colnames(df1)[grepl(".Freedom", colnames(df1))],
                 "FreedomIndex")])
data.frame(cor = cor[-nrow(cor), "FreedomIndex"])
```

##	cor
## Business.Freedom	0.7833618
## Labor.Freedom	0.5722339
## Monetary.Freedom	0.7680385
## Trade.Freedom	0.8125706
## Investment.Freedom	0.8136312
## Financial.Freedom	0.8510606

Rank countries by economic Freedom Index

Top 10 countries

```
df1$FreedomIndexRank <- rank(df1$FreedomIndex)
df1$FreedomIndexRank <- nrow(df1) - df1$FreedomIndexRank + 1
df1 <- df1[order(df1$FreedomIndexRank), ]
rownames(df1) <- NULL

head(df1[, c("Country", "Score", "FreedomIndex", "FreedomIndexRank")], 10)
```

##	Country	Score	FreedomIndex	FreedomIndexRank
## 1	Hong Kong SAR	3.784657	100.00000	1
## 2	Singapore	3.549703	98.36754	2
## 3	Denmark	3.445438	97.64310	3
## 4	New Zealand	3.400626	97.33175	4
## 5	Australia	3.336062	96.88315	5
## 6	United Kingdom	3.051521	94.90616	6
## 7	Switzerland	2.907249	93.90376	7
## 8	Netherlands	2.771262	92.95892	8
## 9	Ireland	2.674441	92.28621	9
## 10	United States	2.670496	92.25879	10

Rank countries by economic Freedom Index

Bottom 10 countries

```
tail(df1[, c("Country", "Score", "FreedomIndex", "FreedomIndexRank")], 10)
```

##	Country	Score	FreedomIndex	FreedomIndexRank
## 171	Congo, Republic of	-2.756953	54.54880	171
## 172	Uzbekistan	-3.048349	52.52417	172
## 173	Afghanistan	-3.088722	52.24366	173
## 174	Sudan	-3.507688	49.33267	174
## 175	Turkmenistan	-3.611616	48.61058	175
## 176	Iran	-3.626171	48.50945	176
## 177	Zimbabwe	-3.629881	48.48367	177
## 178	Cuba	-4.597460	41.76093	178
## 179	Venezuela	-6.101804	31.30872	179
## 180	Korea, North	-10.607942	0.00000	180

t-Distributed Stochastic Neighbor Embedding (t-SNE)



PCA vs t-SNE

PCA is a **linear** feature extraction technique. It performs a linear mapping of the data to a lower-dimensional space in such a way that the variance of the data in the low-dimensional representation is maximized.

t-SNE is a **non-linear** technique for dimensionality reduction that is particularly well suited for the visualization of high-dimensional datasets. It is extensively applied in image processing, NLP, genomic data and speech processing.

t-SNE

The algorithm starts by calculating the probability of similarity of points in high-dimensional space and calculating the probability of similarity of points in the corresponding low-dimensional space. The similarity of points is calculated as the conditional probability that a point i would choose point j as its neighbor if neighbors were picked in proportion to their probability density under a Gaussian (normal distribution) centered at i .

$$p_{j|i} = \frac{\exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma_i^2}\right)}{\sum_{k \neq i} \left(\exp\left(-\frac{\|x_i - x_k\|^2}{2\sigma_k^2}\right)\right)}$$

where σ_i is the variance of the Gaussian that is centered on datapoint x_i . Similar conditional probabilities are defined for the low-dimensional counterparts y_i and y_j of x_i and x_j

$$q_{j|i} = \frac{\exp(-\|y_i - y_j\|^2)}{\sum_{k \neq i} \left(\exp(-\|y_i - y_k\|^2)\right)}$$

Since we are only interested in modeling pairwise similarities, we set $p_{i|i} = q_{i|i} = 0$

t-SNE

It then tries to minimize the difference between these conditional probabilities (or similarities) in higher-dimensional and lower-dimensional space for a perfect representation of data points ($p_{j|i} = q_{j|i} = 0$) in lower-dimensional space (2D or 3D space).

In simpler terms, t-SNE minimizes the divergence between two distributions: a distribution that measures pairwise similarities of the input objects and a distribution that measures pairwise similarities of the corresponding low-dimensional points in the embedding.

In this way, t-SNE maps the multi-dimensional data to a lower dimensional space and attempts to find patterns in the data by identifying observed clusters based on similarity of data points with multiple features.

However!!! after this process, the **input features are no longer identifiable**, and you cannot make any inference based only on the output of t-SNE. Hence it is mainly a **data exploration and visualization technique**.

t-SNE in R


```
#library(Rtsne)
set.seed(1)
tsne <- Rtsne(df1[, -1], dims = 2)
dims <- data.frame(Country = df1$Country, tsne$Y)
head(dims)
```

	Country	X1	X2
## 1	Hong Kong SAR	-12.96044	-0.5297066
## 2	Singapore	-12.71198	-0.1047265
## 3	Denmark	-12.61775	-0.3477270
## 4	New Zealand	-12.47197	-0.1160881
## 5	Australia	-12.64009	-0.6194472
## 6	United Kingdom	-12.05516	-0.9358272

The object **tsne\$Y** contains coordinates of countries in 2D space

Visualizing t-SNE

```
ggplot(dims, aes(x = X1, y = X2, label = Country)) +  
  geom_point() + geom_text(aes(label = Country), hjust = 0, vjust = 0) +  
  xlab("tSNE dimension 1") + ylab("tSNE dimension 2")
```



Thank You!

Questions?